

Assignment 5

March 10, 2019

```
In [1]: # Part 1 (a)
        for line in open('SMSSpamCollection'):
            if line[0] == 's':
                print(line)
                break

        for line in open('SMSSpamCollection'):
            if line[0] == 'h':
                print(line)
                break
```

spam Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 871

ham Go until jurong point, crazy.. Available only in bugis n great world la e buffet...

```
In [2]: # Part 1 (b)
        num_spam = 0
        for line in open('SMSSpamCollection'):
            if line[0] == 's':
                num_spam += 1
        print(num_spam)

        num_non = 0
        for line in open('SMSSpamCollection'):
            if line[0] == 'h':
                num_non += 1
        print(num_non)
```

747
4827

```
In [3]: # Part 1 (c)
        # Advantage: RNN will also learn specific names
        # Disadvantage: computing power and time are spent learning words instead of
        # how words are connected (often more important)
```

```

In [4]: # Part 1 (d)
import torchtext

text_field = torchtext.data.Field(sequential=True,      # text sequence
                                  tokenize=lambda x: x, # because are building a character
                                  include_lengths=True,  # to track the length of sequence
                                  batch_first=True,
                                  use_vocab=True)        # to turn each character into integer
label_field = torchtext.data.Field(sequential=False,   # not a sequence
                                   use_vocab=False,     # don't need to track vocabulary
                                   is_target=True,
                                   batch_first=True,
                                   preprocessing=lambda x: int(x == 'spam')) # convert string to int

fields = [('label', label_field), ('sms', text_field)]
dataset = torchtext.data.TabularDataset("SMSSpamCollection", # name of the file
                                       "tsv",                 # fields are separated by tab
                                       fields)

# dataset[0].sms
# dataset[0].label

train, valid, test= dataset.split(split_ratio = [0.6,0.2,0.2],
                                  stratified = True, strata_field = 'label')

In [5]: len(train)

Out[5]: 3343

In [6]: len(valid)

Out[6]: 1115

In [7]: len(test)

Out[7]: 1114

In [8]: # Part 1 (e)
# A balanced training set will not lead to overfitting to the most popular
# label. For example, if a dataset contains 90% label A, just by outputting A,
# a model can get 90% accuracy.

# save the original training examples
old_train_examples = train.examples
# get all the spam messages in `train`
train_spam = []
for item in train.examples:
    if item.label == 1:
        train_spam.append(item)
# duplicate each spam message 6 more times
train.examples = old_train_examples + train_spam * 6

```

```
In [9]: # Part 1 (f)
        # text_field.vocab.stoi is a string-to-integer dictionary of existing
        # characters to their assigned integer values.
        # text_field.vocab.itos is an integer-to-string dictionary-like list of
        # integer values (represented by the index position) to their characters.
```

```
text_field.build_vocab(train)
#text_field.vocab.stoi
#text_field.vocab.itos
```

```
In [11]: text_field.vocab.stoi
```

```
Out[11]: defaultdict(<function torchtext.vocab._default_unk_index()>,
                    {'<unk>': 0,
                     '<pad>': 1,
                     ' ': 2,
                     'e': 3,
                     'o': 4,
                     't': 5,
                     'a': 6,
                     'n': 7,
                     'r': 8,
                     'i': 9,
                     's': 10,
                     'l': 11,
                     'u': 12,
                     'h': 13,
                     '0': 14,
                     'd': 15,
                     'c': 16,
                     '.': 17,
                     'm': 18,
                     'y': 19,
                     'w': 20,
                     'p': 21,
                     'g': 22,
                     '1': 23,
                     'f': 24,
                     'b': 25,
                     '2': 26,
                     'T': 27,
                     '8': 28,
                     'E': 29,
                     'k': 30,
                     'v': 31,
                     '5': 32,
                     'S': 33,
                     'C': 34,
```

'0': 35,
'I': 36,
'4': 37,
'N': 38,
'A': 39,
'x': 40,
'7': 41,
'6': 42,
'R': 43,
'3': 44,
' ': 45,
'!': 46,
'9': 47,
'P': 48,
'L': 49,
'M': 50,
'U': 51,
'W': 52,
'H': 53,
'D': 54,
'F': 55,
'B': 56,
'Y': 57,
'G': 58,
'/' : 59,
'?' : 60,
'"' : 61,
'č' : 62,
'&' : 63,
'-' : 64,
':' : 65,
'X' : 66,
'V' : 67,
'z' : 68,
'*' : 69,
'j' : 70,
'K' : 71,
')' : 72,
'J' : 73,
';' : 74,
'+' : 75,
'(' : 76,
'Q' : 77,
'q' : 78,
'"' : 79,
'#' : 80,
'>' : 81,
'=' : 82,

```

'@': 83,
'Z': 84,
'ü': 85,
'<': 86,
'$': 87,
'': 88,
'Ü': 89,
'%': 90,
'\x92': 91,
'[': 92,
']': 93,
'|': 94,
'\x93': 95,
'a': 96,
'_': 97,
'': 98,
'': 99,
'': 100,
'ú': 101,
'\t': 102,
'\n': 103,
'\x91': 104,
'\x96': 105,
'\\': 106,
'^': 107,
'~': 108,
'z': 109,
'': 110,
'': 111,
'': 112})

```

In [10]: # Part 1 (g)

```

# <unk> represents an unknown, rare word. Used when dealing at word-level.
# <pad> is padding used to make batches of equal length.

```

In [33]: # Part 1 (h)

```

train_iter = torchtext.data.BucketIterator(train,
                                             batch_size=32,
                                             sort_key=lambda x: len(x.sms), # to minimize
                                             sort_within_batch=True,           # sort within
                                             repeat=False)                       # repeat t

num_pad = []

print('Max length in each batch is the following:')
for i, batch in enumerate(train_iter):
    if i >= 10:

```

```

        break

    print(batch.sms[1][0].item())

    num_pad_batch = 0
    for j in batch.sms[0]:
        for k in j:
            if k.item() == 1:
                num_pad_batch += 1
    num_pad.append(num_pad_batch)

    print('The number of <pad>s in each batch is the following:')
    print(num_pad)

```

Max length in each batch is the following:

```

93
28
160
36
138
146
148
161
41
158

```

The number of <pad>s in each batch is the following:

```

[76, 27, 0, 20, 19, 30, 13, 20, 27, 0]

```

In [13]: *# Part 2*

```

import torch
import numpy

# You mind find this code helpful for obtaining
# pytorch one-hot vectors.

ident = torch.eye(10, dtype = torch.float32)
#ident = torch.eye(10)
print(ident[0]) # one-hot vector
print(ident[1]) # one-hot vector
x = torch.tensor([[1, 2], [3, 4]])
print(ident[x]) # one-hot vectors

```

```

tensor([1., 0., 0., 0., 0., 0., 0., 0., 0., 0.])
tensor([0., 1., 0., 0., 0., 0., 0., 0., 0., 0.])
tensor([[0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
        [0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]])

```

```
[[0., 0., 0., 1., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.]])
```

In [14]: `import torch.nn as nn`

```
class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(RNN, self).__init__()
        self.emb = nn.Embedding(len(text_field.vocab.stoi), hidden_size)
        self.hidden_size = hidden_size
        self.rnn = nn.RNN(input_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, num_classes)
```

```
    def forward(self, x):
        x = self.emb(x)
        h0 = torch.zeros(1, x.size(0), self.hidden_size)
        out, _ = self.rnn(x, h0)
        #out = self.fc(out[:, -1, :])
        out = self.fc(torch.max(out, dim=1)[0])
        return out
```

*# This code is here to help you test your model.
You may need to change this depending on how your forward
function is set up.*

```
model = RNN(10, 10, 2)
sample_batch = next(iter(train_iter))
sms = sample_batch.sms[0]
length = sample_batch.sms[1]
y = model(sms)
print(y.shape)
```

`torch.Size([32, 2])`

In [28]: *# Part 3 (a)*

```
import matplotlib.pyplot as plt
%matplotlib inline
```

```
def get_accuracy(model, data):
    data_iter = torchtext.data.BucketIterator(data,
                                                batch_size=64,
                                                sort_key=lambda x: len(x.sms),
                                                repeat=False)

    correct, total = 0, 0
    for i, batch in enumerate(data_iter):
        output = model(batch.sms[0]) # You may need to modify this, depending on your
```

```

        pred = output.max(1, keepdim=True)[1]
        correct += pred.eq(batch.label.view_as(pred)).sum().item()
        total += batch.sms[1].shape[0]
    return correct / total

def train_rnn_network(model, train, valid, num_epochs=5, learning_rate=1e-5, batch_size=10):
    train_iter = torchtext.data.BucketIterator(train,
                                                batch_size=batch_size,
                                                sort_key=lambda x: len(x.sms), # to minimize padding
                                                sort_within_batch=True, # sort within batch
                                                repeat=False) # repeat training data

    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
    losses, train_acc, valid_acc, loss = [], [], [], 0
    epochs = []
    for epoch in range(num_epochs):
        print(epoch)
        for j, batch in enumerate(train_iter):
            l = batch.label
            s = batch.sms

            optimizer.zero_grad()
            pred = model(s[0])
            loss = criterion(pred, l)

            loss.backward()
            optimizer.step()

        losses.append(float(loss))

        epochs.append(epoch)
        train_acc.append(get_accuracy(RNN(10, 10, 2), train))
        valid_acc.append(get_accuracy(RNN(10, 10, 2), valid))
        print("Epoch %d; Loss %f; Train Acc %f; Val Acc %f" % (
            epoch+1, loss, train_acc[-1], valid_acc[-1]))
    # plotting
    plt.title("Training Curve")
    plt.plot(losses, label="Train")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.show()

    plt.title("Training Curve")
    plt.plot(epochs, train_acc, label="Train")
    plt.plot(epochs, valid_acc, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")

```



```
plt.legend(loc='best')
plt.show()
```

```
In [16]: train_rnn_network(RNN(10, 10, 2), train_iter, valid, num_epochs=5,
                           learning_rate=1e-5)
```

0

Epoch 1; Loss 0.763615; Train Acc 0.480020; Val Acc 0.862780

1

Epoch 2; Loss 0.633558; Train Acc 0.520478; Val Acc 0.864574

2


Epoch 3; Loss 0.606716; Train Acc 0.522799; Val Acc 0.865471


3

Epoch 4; Loss 0.794142; Train Acc 0.519980; Val Acc 0.134529

4

Epoch 5; Loss 0.615960; Train Acc 0.479688; Val Acc 0.141704







```
In [24]: # Part 3 (b)
         # Trial #1
         # Previous num_epochs is 5. Could not see the overall trend well.
         # Increased num_epochs to 20
         # Valid acc is 85% (with consistency problems)

         train_rnn_network(RNN(10, 10, 2), train_iter, valid, num_epochs=20,
                           learning_rate=1e-5)
```

```
0
Epoch 1; Loss 0.888866; Train Acc 0.480020; Val Acc 0.135426
1
Epoch 2; Loss 1.250517; Train Acc 0.520478; Val Acc 0.134529
2
Epoch 3; Loss 1.236499; Train Acc 0.519980; Val Acc 0.863677
3
Epoch 4; Loss 0.345989; Train Acc 0.416846; Val Acc 0.547085
4
Epoch 5; Loss 0.599730; Train Acc 0.519980; Val Acc 0.865471
5
Epoch 6; Loss 1.205293; Train Acc 0.521472; Val Acc 0.134529
6
Epoch 7; Loss 0.452605; Train Acc 0.480020; Val Acc 0.134529
7
```

Epoch 8; Loss 1.155143; Train Acc 0.519980; Val Acc 0.865471
8
Epoch 9; Loss 0.343446; Train Acc 0.519980; Val Acc 0.134529
9
Epoch 10; Loss 0.410172; Train Acc 0.480020; Val Acc 0.865471
10
Epoch 11; Loss 0.595016; Train Acc 0.519980; Val Acc 0.865471
11
Epoch 12; Loss 0.995447; Train Acc 0.519980; Val Acc 0.865471
12
Epoch 13; Loss 0.451761; Train Acc 0.519980; Val Acc 0.134529
13
Epoch 14; Loss 1.029853; Train Acc 0.480020; Val Acc 0.134529
14
Epoch 15; Loss 0.591017; Train Acc 0.480020; Val Acc 0.134529
15
Epoch 16; Loss 1.046984; Train Acc 0.592936; Val Acc 0.134529
16
Epoch 17; Loss 0.508067; Train Acc 0.519980; Val Acc 0.865471
17
Epoch 18; Loss 1.033620; Train Acc 0.438899; Val Acc 0.134529
18
Epoch 19; Loss 0.451655; Train Acc 0.521970; Val Acc 0.843049
19
Epoch 20; Loss 0.848682; Train Acc 0.480020; Val Acc 0.865471







```
In [25]: # Trial #2
         # Previous loss and valid plots very jittery, thought decreasing learning_rate
         # would stablized. It did not.
         # Decreased learning_rate to 1e-6
         # Valid acc is still 85% (with consistency problems)
```

```
train_rnn_network(RNN(10, 10, 2), train_iter, valid, num_epochs=20,
                  learning_rate=1e-6)
```

```
0
Epoch 1; Loss 0.769215; Train Acc 0.524623; Val Acc 0.134529
1
Epoch 2; Loss 0.726124; Train Acc 0.379373; Val Acc 0.865471
2
Epoch 3; Loss 0.608583; Train Acc 0.480020; Val Acc 0.795516
3
Epoch 4; Loss 0.617511; Train Acc 0.586470; Val Acc 0.134529
4
Epoch 5; Loss 0.819547; Train Acc 0.480849; Val Acc 0.135426
5
Epoch 6; Loss 0.628893; Train Acc 0.519980; Val Acc 0.134529
```

6
Epoch 7; Loss 0.621459; Train Acc 0.519980; Val Acc 0.834978
7
Epoch 8; Loss 0.773693; Train Acc 0.413198; Val Acc 0.818834
8
Epoch 9; Loss 0.606781; Train Acc 0.480020; Val Acc 0.134529
9
Epoch 10; Loss 0.627239; Train Acc 0.520146; Val Acc 0.134529
10
Epoch 11; Loss 0.651540; Train Acc 0.480020; Val Acc 0.856502
11
Epoch 12; Loss 0.784925; Train Acc 0.520975; Val Acc 0.134529
12
Epoch 13; Loss 0.667630; Train Acc 0.471398; Val Acc 0.134529
13
Epoch 14; Loss 0.615776; Train Acc 0.480020; Val Acc 0.865471
14
Epoch 15; Loss 0.800693; Train Acc 0.479854; Val Acc 0.865471
15
Epoch 16; Loss 0.618387; Train Acc 0.480020; Val Acc 0.865471
16
Epoch 17; Loss 0.723727; Train Acc 0.480020; Val Acc 0.809865
17
Epoch 18; Loss 0.794773; Train Acc 0.478859; Val Acc 0.860090
18
Epoch 19; Loss 0.801456; Train Acc 0.479688; Val Acc 0.134529
19
Epoch 20; Loss 0.759925; Train Acc 0.519980; Val Acc 0.865471



```

In [26]: # Trial #3
         # Thought a model trained with a balanced dataset would not do well during
         # validation with an imbalanced dataset.
         # Duplicated spam in valid dataset by a factor of 6 also
         # Valid acc is 70% at peek

old_valid_examples = valid.examples
valid_spam = []
for item in valid.examples:
    if item.label == 1:
        valid_spam.append(item)


valid.examples = old_valid_examples + valid_spam * 6


train_rnn_network(RNN(10, 10, 2), train_iter, valid, num_epochs=20,
                  learning_rate=1e-6)

0
Epoch 1; Loss 0.641894; Train Acc 0.521970; Val Acc 0.478908
1
Epoch 2; Loss 0.666710; Train Acc 0.479854; Val Acc 0.477916
2
Epoch 3; Loss 0.822245; Train Acc 0.480020; Val Acc 0.478908
3
Epoch 4; Loss 0.808926; Train Acc 0.480020; Val Acc 0.521092
4
Epoch 5; Loss 0.828313; Train Acc 0.519980; Val Acc 0.500744
5
Epoch 6; Loss 0.817445; Train Acc 0.480020; Val Acc 0.524069
6
Epoch 7; Loss 0.702273; Train Acc 0.519980; Val Acc 0.522084
7
Epoch 8; Loss 0.808636; Train Acc 0.529100; Val Acc 0.521092
8
Epoch 9; Loss 0.649615; Train Acc 0.480020; Val Acc 0.521092
9
Epoch 10; Loss 0.648863; Train Acc 0.480020; Val Acc 0.478908
10
Epoch 11; Loss 0.848926; Train Acc 0.514508; Val Acc 0.533995
11
Epoch 12; Loss 0.842385; Train Acc 0.519980; Val Acc 0.522084
12
Epoch 13; Loss 0.854700; Train Acc 0.479688; Val Acc 0.478908
13
Epoch 14; Loss 0.829356; Train Acc 0.480020; Val Acc 0.478908
14
Epoch 15; Loss 0.695435; Train Acc 0.480020; Val Acc 0.478412
15

```

Epoch 16; Loss 0.793028; Train Acc 0.480352; Val Acc 0.478412
16
Epoch 17; Loss 0.582858; Train Acc 0.521970; Val Acc 0.697767
17
Epoch 18; Loss 0.586601; Train Acc 0.481844; Val Acc 0.478908
18
Epoch 19; Loss 0.734364; Train Acc 0.519980; Val Acc 0.529032
19
Epoch 20; Loss 0.574606; Train Acc 0.480020; Val Acc 0.521092







```
In [29]: # Trial #4
# Thought previous plot was too stable, decreasing batch_size may bring more
# improvement.
# Decreased batch_size to 8
# Valid acc is 63% at peek
```

```
train_rnn_network(RNN(10, 10, 2), train, valid, num_epochs=20,
                  learning_rate=1e-6, batch_size=8)
```

```
0
Epoch 1; Loss 0.506941; Train Acc 0.480020; Val Acc 0.521092
1
Epoch 2; Loss 0.919603; Train Acc 0.519980; Val Acc 0.522084
2
Epoch 3; Loss 0.732741; Train Acc 0.480020; Val Acc 0.478908
3
Epoch 4; Loss 0.904622; Train Acc 0.520312; Val Acc 0.526551
4
Epoch 5; Loss 0.847900; Train Acc 0.519980; Val Acc 0.478412
5
Epoch 6; Loss 0.787778; Train Acc 0.478528; Val Acc 0.623821
6
Epoch 7; Loss 0.516550; Train Acc 0.476538; Val Acc 0.521092
7
```

Epoch 8; Loss 0.877684; Train Acc 0.480020; Val Acc 0.475931
8
Epoch 9; Loss 0.847396; Train Acc 0.520643; Val Acc 0.390074
9
Epoch 10; Loss 0.512623; Train Acc 0.480020; Val Acc 0.521092
10
Epoch 11; Loss 0.835399; Train Acc 0.523130; Val Acc 0.475931
11
Epoch 12; Loss 0.562757; Train Acc 0.661416; Val Acc 0.478908
12
Epoch 13; Loss 0.667825; Train Acc 0.702703; Val Acc 0.478412
13
Epoch 14; Loss 0.677824; Train Acc 0.519980; Val Acc 0.478908
14
Epoch 15; Loss 0.712267; Train Acc 0.519980; Val Acc 0.474938
15
Epoch 16; Loss 0.906640; Train Acc 0.519980; Val Acc 0.478908
16
Epoch 17; Loss 0.883359; Train Acc 0.489139; Val Acc 0.478908
17
Epoch 18; Loss 0.573429; Train Acc 0.519980; Val Acc 0.478908
18
Epoch 19; Loss 0.569902; Train Acc 0.480020; Val Acc 0.478908
19
Epoch 20; Loss 0.566561; Train Acc 0.480020; Val Acc 0.521588




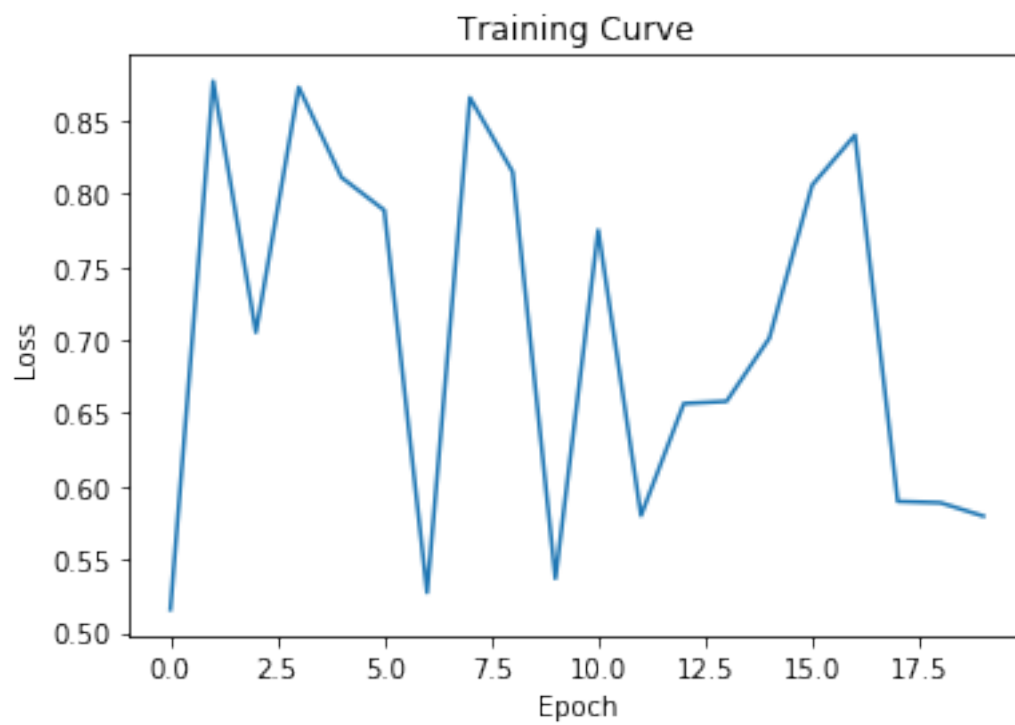


```
In [30]: # Trial #5
         # Previous hyper-parameters ran reasonably quick, thought increasing layer
         # sizes could improve accuracy.
         # Increased input and hidden layer sizes to 20
         # Valid acc is 63% at peek
```

```
train_rnn_network(RNN(20, 20, 2), train, valid, num_epochs=20,
                  learning_rate=1e-6, batch_size=8)
```

```
0
Epoch 1; Loss 0.515827; Train Acc 0.519980; Val Acc 0.521092
1
Epoch 2; Loss 0.877051; Train Acc 0.519980; Val Acc 0.478908
2
Epoch 3; Loss 0.705355; Train Acc 0.519980; Val Acc 0.478908
3
Epoch 4; Loss 0.872909; Train Acc 0.480020; Val Acc 0.521092
4
Epoch 5; Loss 0.811247; Train Acc 0.583320; Val Acc 0.478908
5
Epoch 6; Loss 0.788810; Train Acc 0.480020; Val Acc 0.521092
```

6
Epoch 7; Loss 0.527461; Train Acc 0.480020; Val Acc 0.478908
7
Epoch 8; Loss 0.865794; Train Acc 0.520643; Val Acc 0.521092
8
Epoch 9; Loss 0.815014; Train Acc 0.521141; Val Acc 0.521092
9
Epoch 10; Loss 0.537031; Train Acc 0.570552; Val Acc 0.478908
10
Epoch 11; Loss 0.775206; Train Acc 0.369093; Val Acc 0.468486
11
Epoch 12; Loss 0.580069; Train Acc 0.480020; Val Acc 0.474938
12
Epoch 13; Loss 0.656616; Train Acc 0.519980; Val Acc 0.530521
13
Epoch 14; Loss 0.658153; Train Acc 0.478030; Val Acc 0.478908
14
Epoch 15; Loss 0.701381; Train Acc 0.480020; Val Acc 0.478908
15
Epoch 16; Loss 0.805945; Train Acc 0.519980; Val Acc 0.521092
16
Epoch 17; Loss 0.840178; Train Acc 0.535898; Val Acc 0.521092
17
Epoch 18; Loss 0.589724; Train Acc 0.479357; Val Acc 0.426303
18
Epoch 19; Loss 0.588810; Train Acc 0.477533; Val Acc 0.629280
19
Epoch 20; Loss 0.579693; Train Acc 0.576190; Val Acc 0.478412



```



In [31]: # Trial #6
         # Previous plots' shapes too unorthodox, thought increasing batch_size would
         # bring more stability and more traditional plot shape.
         # Increased batch_size to 256
         # Valid acc is 58% at peek

         train_rnn_network(RNN(20, 20, 2), train, valid, num_epochs=20,
                           learning_rate=1e-6, batch_size=256)

0
Epoch 1; Loss 1.076035; Train Acc 0.519980; Val Acc 0.351861
1
Epoch 2; Loss 0.466367; Train Acc 0.519980; Val Acc 0.477916
2
Epoch 3; Loss 0.728957; Train Acc 0.474714; Val Acc 0.521588
3
Epoch 4; Loss 1.140865; Train Acc 0.479191; Val Acc 0.478908
4
Epoch 5; Loss 1.068177; Train Acc 0.479357; Val Acc 0.521588
5
Epoch 6; Loss 0.556215; Train Acc 0.480020; Val Acc 0.478908
6
Epoch 7; Loss 1.067514; Train Acc 0.519980; Val Acc 0.521092
7
Epoch 8; Loss 1.164998; Train Acc 0.478196; Val Acc 0.478908
8
Epoch 9; Loss 0.553509; Train Acc 0.519980; Val Acc 0.521092
9
Epoch 10; Loss 0.433946; Train Acc 0.476870; Val Acc 0.521092
10
Epoch 11; Loss 1.127679; Train Acc 0.480020; Val Acc 0.521092
11
Epoch 12; Loss 0.953405; Train Acc 0.682806; Val Acc 0.478908
12
Epoch 13; Loss 0.971642; Train Acc 0.480020; Val Acc 0.478908
13
Epoch 14; Loss 0.424787; Train Acc 0.480020; Val Acc 0.467990
14
Epoch 15; Loss 0.424081; Train Acc 0.621290; Val Acc 0.521092
15
Epoch 16; Loss 0.958184; Train Acc 0.480020; Val Acc 0.478908
16
Epoch 17; Loss 0.474263; Train Acc 0.478030; Val Acc 0.569727
17
Epoch 18; Loss 0.432907; Train Acc 0.519980; Val Acc 0.478908
18
Epoch 19; Loss 1.126798; Train Acc 0.480020; Val Acc 0.521092
19

```

Epoch 20; Loss 0.407969; Train Acc 0.480020; Val Acc 0.521092



```
In [32]: # Trial #7
         # Previous change worked somewhat well, thought further increasing the
         # batch_size would improve result like last time. It did not work.
         # Increased batch_size to 1024
         # Valid acc is 53% at peek
```

```
train_rnn_network(RNN(20, 20, 2), train, valid, num_epochs=20,
                  learning_rate=1e-6, batch_size=1024)
```


```
0
Epoch 1; Loss 0.949650; Train Acc 0.519980; Val Acc 0.521092
1
Epoch 2; Loss 0.697485; Train Acc 0.480020; Val Acc 0.455583
2
Epoch 3; Loss 0.922858; Train Acc 0.519980; Val Acc 0.478908
3
Epoch 4; Loss 0.469681; Train Acc 0.480020; Val Acc 0.521092
4
Epoch 5; Loss 0.951101; Train Acc 0.519980; Val Acc 0.521092
5
Epoch 6; Loss 0.696137; Train Acc 0.480020; Val Acc 0.452109
6
Epoch 7; Loss 0.697391; Train Acc 0.523628; Val Acc 0.478908
7
Epoch 8; Loss 0.694923; Train Acc 0.519980; Val Acc 0.531514
8
Epoch 9; Loss 0.951492; Train Acc 0.519980; Val Acc 0.521092
9
Epoch 10; Loss 0.921834; Train Acc 0.480020; Val Acc 0.478908
10
Epoch 11; Loss 0.981023; Train Acc 0.471895; Val Acc 0.521092
11
Epoch 12; Loss 0.922342; Train Acc 0.519980; Val Acc 0.478908
12
Epoch 13; Loss 0.470069; Train Acc 0.521141; Val Acc 0.472953
13
Epoch 14; Loss 0.695546; Train Acc 0.519980; Val Acc 0.418362
14
Epoch 15; Loss 0.980040; Train Acc 0.524125; Val Acc 0.521092
15
Epoch 16; Loss 0.922818; Train Acc 0.480020; Val Acc 0.522581
16
Epoch 17; Loss 0.470106; Train Acc 0.519980; Val Acc 0.521092
17
Epoch 18; Loss 0.951533; Train Acc 0.474382; Val Acc 0.478908
```



18

Epoch 19; Loss 0.694851; Train Acc 0.479854; Val Acc 0.316129

19

Epoch 20; Loss 0.978571; Train Acc 0.478362; Val Acc 0.447146





```
In [ ]: # Part 3 (c)
        # Best result is trial #5. Valid acc is 63% with decent overall shape.
        # More complex model likely will help accuracy. Current one quite simple.

In [ ]: # Part 4
        # Not too difficult task.
        # A baseline model could be a word-level RNN looking for keywords to detect spam.
        # In my experience, spam emails are not written creatively. Thus, detecting
        # a few keywords, such as "free", "lottery", "gift", "trial", is likely to
        # give a decent result.
```